

# Machine Learning Engineer Nanodegree

## Model Evaluation & Validation

### Project: Predicting Boston Housing Prices

Welcome to the first project of the Machine Learning Engineer Nanodegree! In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**Implementation**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

**Note:** Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. In addition, Markdown cells can be edited by typically double-clicking the cell to enter edit mode.

## Getting Started

In this project, you will evaluate the performance and predictive power of a model that has been trained and tested on data collected from homes in suburbs of Boston, Massachusetts. A model trained on this data that is seen as a *good fit* could then be used to make certain predictions about a home — in particular, its monetary value. This model would prove to be invaluable for someone like a real estate agent who could make use of such information on a daily basis.

The dataset for this project originates from the [UCI Machine Learning Repository](https://archive.ics.uci.edu/ml/machine-learning-databases/housing/) (<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/>). The Boston housing data was collected in 1978 and each of the 506 entries represent aggregated data about 14 features for homes from various suburbs in Boston, Massachusetts. For the purposes of this project, the following preprocessing steps have been made to the dataset:

- 16 data points have an 'MEDV' value of 50.0. These data points likely contain **missing or censored values** and have been removed.
- 1 data point has an 'RM' value of 8.78. This data point can be considered an **outlier** and has been removed.
- The features 'RM', 'LSTAT', 'PTRATIO', and 'MEDV' are essential. The remaining **non-relevant features** have been excluded.
- The feature 'MEDV' has been **multiplicatively scaled** to account for 35 years of market inflation.

Run the code cell below to load the Boston housing dataset, along with a few of the necessary Python libraries required for this project. You will know the dataset loaded successfully if the size of the dataset is reported.

```
In [2]: # Import libraries necessary for this project
import numpy as np
import pandas as pd
from sklearn.cross_validation import ShuffleSplit

# Import supplementary visualizations code visuals.py
import visuals as vs

# Pretty display for notebooks
%matplotlib inline

# Load the Boston housing dataset
data = pd.read_csv('housing.csv')
prices = data['MEDV']
features = data.drop('MEDV', axis = 1)

# Success
print("Boston housing dataset has {} data points with {} variables each."
      .format(*data.shape))

print(data.head())
```

Boston housing dataset has 489 data points with 4 variables each.

	RM	LSTAT	PTRATIO	MEDV
0	6.575	4.98	15.3	504000.0
1	6.421	9.14	17.8	453600.0
2	7.185	4.03	17.8	728700.0
3	6.998	2.94	18.7	701400.0
4	7.147	5.33	18.7	760200.0

## Data Exploration

In this first section of this project, you will make a cursory investigation about the Boston housing data and provide your observations. Familiarizing yourself with the data through an explorative process is a fundamental practice to help you better understand and justify your results.

Since the main goal of this project is to construct a working model which has the capability of predicting the value of houses, we will need to separate the dataset into **features** and the **target variable**. The **features**, 'RM' , 'LSTAT' , and 'PTRATIO' , give us quantitative information about each data point. The **target variable**, 'MEDV' , will be the variable we seek to predict. These are stored in `features` and `prices` , respectively.

## Implementation: Calculate Statistics

For your very first coding implementation, you will calculate descriptive statistics about the Boston housing prices. Since `numpy` has already been imported for you, use this library to perform the necessary calculations. These statistics will be extremely important later on to analyze various prediction results from the constructed model.

In the code cell below, you will need to implement the following:

- Calculate the minimum, maximum, mean, median, and standard deviation of `'MEDV'`, which is stored in `prices`.
  - Store each calculation in their respective variable.

```
In [3]: # TODO: Minimum price of the data
minimum_price = np.min(prices)

# TODO: Maximum price of the data
maximum_price = np.max(prices)

# TODO: Mean price of the data
mean_price = np.mean(prices)

# TODO: Median price of the data
median_price = np.median(prices)

# TODO: Standard deviation of prices of the data
std_price = np.std(prices)

# Show the calculated statistics
print("Statistics for Boston housing dataset:\n")
print("Minimum price: ${}".format(minimum_price))
print("Maximum price: ${}".format(maximum_price))
print("Mean price: ${}".format(mean_price))
print("Median price ${}".format(median_price))
print("Standard deviation of prices: ${}".format(std_price))
```

Statistics for Boston housing dataset:

Minimum price: \$105000.0  
Maximum price: \$1024800.0  
Mean price: \$454342.9447852761  
Median price \$438900.0  
Standard deviation of prices: \$165171.13154429474

## Question 1 - Feature Observation

As a reminder, we are using three features from the Boston housing dataset: 'RM' , 'LSTAT' , and 'PTRATIO' . For each data point (neighborhood):

- 'RM' is the average number of rooms among homes in the neighborhood.
- 'LSTAT' is the percentage of homeowners in the neighborhood considered "lower class" (working poor).
- 'PTRATIO' is the ratio of students to teachers in primary and secondary schools in the neighborhood.

**Using your intuition, for each of the three features above, do you think that an increase in the value of that feature would lead to an increase in the value of 'MEDV' or a decrease in the value of 'MEDV' ? Justify your answer for each.**

**Hint:** This problem can be phrased using examples like below.

- Would you expect a home that has an 'RM' value(number of rooms) of 6 be worth more or less than a home that has an 'RM' value of 7?
- Would you expect a neighborhood that has an 'LSTAT' value(percent of lower class workers) of 15 have home prices be worth more or less than a neighborhood that has an 'LSTAT' value of 20?
- Would you expect a neighborhood that has an 'PTRATIO' value(ratio of students to teachers) of 10 have home prices be worth more or less than a neighborhood that has an 'PTRATIO' value of 15?

## Answer:

We can say that with increase in RM (no of rooms) the MEDV or price of the house will increase. Here we should also consider the outliers, i.e. a house with 7 rooms and poor built and living condition can be cheaper than a house with 4 or 5 rooms but excellent construction, facilities and maintenance.

For LSTAT: We can say that the LSTAT means  $x\%$  of poor families living in the neighborhood, i.e. that  $x\%$  people living there were low income families but they could still afford the home which can drive the conclusion that if  $y < x$ , i.e. a neighborhood with  $y\%$  ( $y < x$ ) of low income families live there, then the housing is less affordable (or costlier than  $x$ ) than  $x$  neighborhood.

And with PTRATIO we can say that lesser the PTRATIO means less no of students per teacher and that means that education will be good and when education is good in a neighborhood, housing prices also elevate since it will be more costlier (in most cases).

---

## Developing a Model

In this second section of the project, you will develop the tools and techniques necessary for a model to make a prediction. Being able to make accurate evaluations of each model's performance through the use of these tools and techniques helps to greatly reinforce the confidence in your predictions.

## Implementation: Define a Performance Metric

It is difficult to measure the quality of a given model without quantifying its performance over training and testing. This is typically done using some type of performance metric, whether it is through calculating some type of error, the goodness of fit, or some other useful measurement. For this project, you will be calculating the [coefficient of determination](http://stattrek.com/statistics/dictionary.aspx?definition=coefficient_of_determination) ([http://stattrek.com/statistics/dictionary.aspx?definition=coefficient\\_of\\_determination](http://stattrek.com/statistics/dictionary.aspx?definition=coefficient_of_determination)),  $R^2$ , to quantify your model's performance. The coefficient of determination for a model is a useful statistic in regression analysis, as it often describes how "good" that model is at making predictions.

The values for  $R^2$  range from 0 to 1, which captures the percentage of squared correlation between the predicted and actual values of the **target variable**. A model with an  $R^2$  of 0 is no better than a model that always predicts the *mean* of the target variable, whereas a model with an  $R^2$  of 1 perfectly predicts the target variable. Any value between 0 and 1 indicates what percentage of the target variable, using this model, can be explained by the **features**. *A model can be given a negative  $R^2$  as well, which indicates that the model is arbitrarily worse than one that always predicts the mean of the target variable.*

For the `performance_metric` function in the code cell below, you will need to implement the following:

- Use `r2_score` from `sklearn.metrics` to perform a performance calculation between `y_true` and `y_predict`.
- Assign the performance score to the `score` variable.

```
In [4]: # TODO: Import 'r2_score'
        from sklearn.metrics import r2_score

        def performance_metric(y_true, y_predict):
            """ Calculates and returns the performance score between
                true and predicted values based on the metric chosen. """

            # TODO: Calculate the performance score between 'y_true' and 'y_predict'
            score = r2_score( y_true, y_predict)

            # Return the score
            return score
```

## Question 2 - Goodness of Fit

Assume that a dataset contains five data points and a model made the following predictions for the target variable:

True Value	Prediction
3.0	2.5
-0.5	0.0
2.0	2.1
7.0	7.8
4.2	5.3

Run the code cell below to use the `performance_metric` function and calculate this model's coefficient of determination.

```
In [5]: # Calculate the performance of this model
score = performance_metric([3, -0.5, 2, 7, 4.2], [2.5, 0.0, 2.1, 7.8, 5.3])
print("Model has a coefficient of determination, R^2, of {:.3f}.".format(score))
```

Model has a coefficient of determination, R<sup>2</sup>, of 0.923.

- Would you consider this model to have successfully captured the variation of the target variable?
- Why or why not?

**Hint:** The R<sup>2</sup> score is the proportion of the variance in the dependent variable that is predictable from the independent variable. In other words:

- R<sup>2</sup> score of 0 means that the dependent variable cannot be predicted from the independent variable.
- R<sup>2</sup> score of 1 means the dependent variable can be predicted from the independent variable.
- R<sup>2</sup> score between 0 and 1 indicates the extent to which the dependent variable is predictable. An
- R<sup>2</sup> score of 0.40 means that 40 percent of the variance in Y is predictable from X.

**Answer:**

The Basic definition of  $R^2$  can be given as the proportion of the variance in the dependent variable that is predictable from the independent variable(s). The  $R^2$ , coefficient of determination is a statistical measure of how well the regression predictions approximate the real data points. An  $R^2$  of 1 indicates that the regression predictions perfectly fit the data.

Yes I would consider this model to have successfully captured the variation of the target variable since the  $R^2$  score is 0.93, which is very high and we know that a model with  $R^2$  score of 1 is considered to be a model that can perfectly predict the target variable from the independent variables.

**Implementation: Shuffle and Split Data**

Your next implementation requires that you take the Boston housing dataset and split the data into training and testing subsets. Typically, the data is also shuffled into a random order when creating the training and testing subsets to remove any bias in the ordering of the dataset.

For the code cell below, you will need to implement the following:

- Use `train_test_split` from `sklearn.cross_validation` to shuffle and split the `features` and `prices` data into training and testing sets.
  - Split the data into 80% training and 20% testing.
  - Set the `random_state` for `train_test_split` to a value of your choice. This ensures results are consistent.
- Assign the train and testing splits to `x_train`, `x_test`, `y_train`, and `y_test`.

```
In [10]: # TODO: Import 'train_test_split'
from sklearn.cross_validation import train_test_split
import random
random.seed(42)

# TODO: Shuffle and split the data into training and testing subsets
X_train, X_test, y_train, y_test = train_test_split(features, prices, te
st_size=0.2, random_state=42)

# Success
print("Training and testing split was successful.")
```

Training and testing split was successful.



## Question 3 - Training and Testing

- What is the benefit to splitting a dataset into some ratio of training and testing subsets for a learning algorithm?

**Hint:** Think about how overfitting or underfitting is contingent upon how splits on data is done.

## Answer: \*\*

We split the data into training set to train our model/learning algorithm and into testing set to test how our model functions. When we don't do this and instead use the whole data to train our model, the model happens to memorize about the particular data set given instead of learning about the characteristics of the data given, so it becomes necessary to have testing set. Talking about the ratio of split (i.e. 80:20 in our case), the split should be such that not much of information is left out nor all of the information is taken. That is why we seem to train our model with 80% of the data and then use the 20% of data (testing data) to test our model.

---

## Analyzing Model Performance


In this third section of the project, you'll take a look at several models' learning and testing performances on various subsets of training data. Additionally, you'll investigate one particular algorithm with an increasing `'max_depth'` parameter on the full training set to observe how model complexity affects performance. Graphing your model's performance based on varying criteria can be beneficial in the analysis process, such as visualizing behavior that may not have been apparent from the results alone.

## Learning Curves

The following code cell produces four graphs for a decision tree model with different maximum depths. Each graph visualizes the learning curves of the model for both training and testing as the size of the training set is increased. Note that the shaded region of a learning curve denotes the uncertainty of that curve (measured as the standard deviation). The model is scored on both the training and testing sets using  $R^2$ , the coefficient of determination.

Run the code cell below and use these graphs to answer the following question.

```
In [11]: # Produce learning curves for varying training set sizes and maximum depths
vs.ModelLearning(features, prices)
```



## Question 4 - Learning the Data

- Choose one of the graphs above and state the maximum depth for the model.
- What happens to the score of the training curve as more training points are added? What about the testing curve?
- Would having more training points benefit the model?

**Hint:** Are the learning curves converging to particular scores? Generally speaking, the more data you have, the better. But if your training and testing curves are converging with a score above your benchmark threshold, would this be necessary? Think about the pros and cons of adding more training points based on if the training and testing curves are converging.

## Answer:

The graph that I would select from above is the graph on upper right corner with `max_depth=3`.

As more and more training points are added we will lean towards overfitting the model. the training curve will tend to stay near 1 since we are scoring based on R2 score. And the testing curve will tend to converge with the training curve.


No, it seems that adding more and more training points will benefit the model but the case is adding more and more no of training points will make the model memorize the data set instead of learning about its features.

## Complexity Curves

The following code cell produces a graph for a decision tree model that has been trained and validated on the training data using different maximum depths. The graph produces two complexity curves — one for training and one for validation. Similar to the **learning curves**, the shaded regions of both the complexity curves denote the uncertainty in those curves, and the model is scored on both the training and validation sets using the `performance_metric` function.

Run the code cell below and use this graph to answer the following two questions Q5 and Q6.

```
In [12]: vs.ModelComplexity(X_train, y_train)
```



### Question 5 - Bias-Variance Tradeoff

- When the model is trained with a maximum depth of 1, does the model suffer from high bias or from high variance?
- How about when the model is trained with a maximum depth of 10? What visual cues in the graph justify your conclusions?

**Hint:** High bias is a sign of underfitting(model is not complex enough to pick up the nuances in the data) and high variance is a sign of overfitting(model is by-hearting the data and cannot generalize well). Think about which model(depth 1 or 10) aligns with which part of the tradeoff.

**Answer:**

When the model is trained with a maximum depth of 1, the model suffers from high bias.

When the model is trained with a max\_depth of 10, the model suffers from High-Variance. We can see from the above plot that the validation score decreases, and both the scores start to unconverge or move apart.

**Question 6 - Best-Guess Optimal Model**

- Which maximum depth do you think results in a model that best generalizes to unseen data?
- What intuition led you to this answer?

**Hint:** Look at the graph above Question 5 and see where the validation scores lie for the various depths that have been assigned to the model. Does it get better with increased depth? At what point do we get our best validation score without overcomplicating our model? And remember, Occam's Razor states "Among competing hypotheses, the one with the fewest assumptions should be selected."

**Answer:**

The maximum depth that I think best generalizes to unseen data is max\_depth=4.

The fact that from given depths, the validation score goes on to decrease after depth=4, and also the fact that both the scores tend to unconverge or move apart is main reason behind selecting the maximum depth of model = 4.

---

**Evaluating Model Performance**

In this final section of the project, you will construct a model and make a prediction on the client's feature set using an optimized model from `fit_model`.

**Question 7 - Grid Search**

- What is the grid search technique?
- How it can be applied to optimize a learning algorithm?

**Hint:** When explaining the Grid Search technique, be sure to touch upon why it is used, what the 'grid' entails and what the end goal of this method is. To solidify your answer, you can also give an example of a parameter in a model that can be optimized using this approach.

**\*\*Answer:**

Grid Search: In Machine Learning we pick up best combination of parameters and hyper-parameters for our model. When we have simple model with 1 hyper-parameter, it is simple to pick up the best model. But when there are more than 1 hyper-parameters, and we need to select a combinations from those, that is where Grid Search comes into play.

In grid search we make table with all the possibilities of the hyper-parameter, and then select the best combinations for our model.

Eg: When we have to train a support vector machine, we have 2 paramters namely: Kernel and C (gamma parameter). The kernel can take up values including linear or polynomial. The gamma parameter can take up values including 0.1, 1, 10 etc So we have various options to choose from and here grid search can help us out by making a table and selecting the best combination of parameters.

**Question 8 - Cross-Validation**

- What is the k-fold cross-validation training technique?
- What benefit does this technique provide for grid search when optimizing a model?

**Hint:** When explaining the k-fold cross validation technique, be sure to touch upon what 'k' is, how the dataset is split into different parts for training and testing and the number of times it is run based on the 'k' value.

When thinking about how k-fold cross validation helps grid search, think about the main drawbacks of grid search which are hinged upon **using a particular subset of data for training or testing** and how k-fold cv could help alleviate that. You can refer to the [docs \(http://scikit-learn.org/stable/modules/cross\\_validation.html#cross-validation\)](http://scikit-learn.org/stable/modules/cross_validation.html#cross-validation) for your answer.

**Answer:**

When we split our data into training and testing set, we tend to throw away some info ( that could potentially be used to train our model) by not using the testing set. The solution to this problem is K-Fold Cross-Validation technique in which, we break our data into k-buckets. Then, we just train our model 'K' times, each time using a different bucket as our testing set and the remaining points/buckets as our training set. Then we take average of the evaluation metric of all the models to get unbiased estimate of model generation on unseen data.

For Example: Lets say we have a sample data of 100 rows. We take 20% as testing set and keep the 20 rows aside. And from the 80 rows of data we decide to divide it into 10 folds, which makes our k =10 and one fold consist of 8 rows. So now we take 1 fold and keep it aside for validation and perform training with the rest 9 folds or 72 rows of data. These process is repeated for each fold and a different validating set everytime.

**<Dear reviewer, help me out on this part. I am not able to relate K-fold and Grid Search**

The benefits of the K-fold technique is that we dont miss out on any particular subset of our data to train and test the model.

## Implementation: Fitting a Model

Your final implementation requires that you bring everything together and train a model using the **decision tree algorithm**. To ensure that you are producing an optimized model, you will train the model using the grid search technique to optimize the 'max\_depth' parameter for the decision tree. The 'max\_depth' parameter can be thought of as how many questions the decision tree algorithm is allowed to ask about the data before making a prediction. Decision trees are part of a class of algorithms called *supervised learning algorithms*.

In addition, you will find your implementation is using `ShuffleSplit()` for an alternative form of cross-validation (see the 'cv\_sets' variable). While it is not the K-Fold cross-validation technique you describe in **Question 8**, this type of cross-validation technique is just as useful!. The `ShuffleSplit()` implementation below will create 10 ( 'n\_splits' ) shuffled sets, and for each shuffle, 20% ( 'test\_size' ) of the data will be used as the *validation set*. While you're working on your implementation, think about the contrasts and similarities it has to the K-fold cross-validation technique.

Please note that `ShuffleSplit` has different parameters in scikit-learn versions 0.17 and 0.18. For the `fit_model` function in the code cell below, you will need to implement the following:

- Use `DecisionTreeRegressor` (<http://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeRegressor.html>) from `sklearn.tree` to create a decision tree regressor object.
  - Assign this object to the 'regressor' variable.
- Create a dictionary for 'max\_depth' with the values from 1 to 10, and assign this to the 'params' variable.
- Use `make_scorer` ([http://scikit-learn.org/stable/modules/generated/sklearn.metrics.make\\_scorer.html](http://scikit-learn.org/stable/modules/generated/sklearn.metrics.make_scorer.html)) from `sklearn.metrics` to create a scoring function object.
  - Pass the `performance_metric` function as a parameter to the object.
  - Assign this scoring function to the 'scoring\_fnc' variable.
- Use `GridSearchCV` ([http://scikit-learn.org/0.17/modules/generated/sklearn.grid\\_search.GridSearchCV.html](http://scikit-learn.org/0.17/modules/generated/sklearn.grid_search.GridSearchCV.html)) from `sklearn.grid_search` to create a grid search object.
  - Pass the variables 'regressor', 'params', 'scoring\_fnc', and 'cv\_sets' as parameters to the object.
  - Assign the `GridSearchCV` object to the 'grid' variable.

```

In [13]: # TODO: Import 'make_scorer', 'DecisionTreeRegressor', and 'GridSearchCV'
from sklearn.metrics import make_scorer
from sklearn.model_selection import GridSearchCV
from sklearn.tree import DecisionTreeRegressor

def fit_model(X, y):
    """ Performs grid search over the 'max_depth' parameter for a
        decision tree regressor trained on the input data [X, y]. """

    # Create cross-validation sets from the training data
    # sklearn version 0.18: ShuffleSplit(n_splits=10, test_size=0.1, train_size=None, random_state=None)
    # sklearn version 0.17: ShuffleSplit(n, n_iter=10, test_size=0.1, train_size=None, random_state=None)
    cv_sets = ShuffleSplit(X.shape[0], n_iter = 10, test_size = 0.20, random_state = 0)

    # TODO: Create a decision tree regressor object
    regressor = DecisionTreeRegressor()

    # TODO: Create a dictionary for the parameter 'max_depth' with a range from 1 to 10
    params = {'max_depth': range(1,11)}

    # TODO: Transform 'performance_metric' into a scoring function using 'make_scorer'
    scoring_fnc = make_scorer(performance_metric)

    # TODO: Create the grid search cv object --> GridSearchCV()
    # Make sure to include the right parameters in the object:
    # (estimator, param_grid, scoring, cv) which have values 'regressor', 'params', 'scoring_fnc', and 'cv_sets' respectively.
    grid = GridSearchCV(regressor, param_grid = params, scoring = scoring_fnc, cv = cv_sets)

    # Fit the grid search object to the data to compute the optimal model
    grid = grid.fit(X, y)

    # Return the optimal model after fitting the data
    return grid.best_estimator_

```

## Making Predictions

Once a model has been trained on a given set of data, it can now be used to make predictions on new sets of input data. In the case of a *decision tree regressor*, the model has learned *what the best questions to ask about the input data are*, and can respond with a prediction for the **target variable**. You can use these predictions to gain information about data where the value of the target variable is unknown — such as data the model was not trained on.

## Question 9 - Optimal Model

- What maximum depth does the optimal model have? How does this result compare to your guess in Question 6?

Run the code block below to fit the decision tree regressor to the training data and produce an optimal model.

```
In [14]: # Fit the training data to the model using grid search
reg = fit_model(X_train, y_train)

# Produce the value for 'max_depth'
print("Parameter 'max_depth' is {} for the optimal model.".format(reg.get_params()['max_depth']))
```

Parameter 'max\_depth' is 4 for the optimal model.

**Hint:** The answer comes from the output of the code snippet above.

**Answer:** max\_depth = 4

The answer over here matches with the answer mentioned in Q6. And we can say that our initial assumption of the Training Error and CV Error converging at highest level with maximum amount of training points comes out to be true and here the model predicts the same.

## Question 10 - Predicting Selling Prices

Imagine that you were a real estate agent in the Boston area looking to use this model to help price homes owned by your clients that they wish to sell. You have collected the following information from three of your clients:

Feature	Client 1	Client 2	Client 3
Total number of rooms in home	5 rooms	4 rooms	8 rooms
Neighborhood poverty level (as %)	17%	32%	3%
Student-teacher ratio of nearby schools	15-to-1	22-to-1	12-to-1

- What price would you recommend each client sell his/her home at?
- Do these prices seem reasonable given the values for the respective features?

**Hint:** Use the statistics you calculated in the **Data Exploration** section to help justify your response. Of the three clients, client 3 has the biggest house, in the best public school neighborhood with the lowest poverty level; while client 2 has the smallest house, in a neighborhood with a relatively high poverty rate and not the best public schools.

Run the code block below to have your optimized model make predictions for each client's home.



```
In [15]: # Produce a matrix for client data
client_data = [[5, 17, 15], # Client 1
               [4, 32, 22], # Client 2
               [8, 3, 12]] # Client 3

# Show predictions
for i, price in enumerate(reg.predict(client_data)):
    print("Predicted selling price for Client {}'s home: ${:,.2f}".format(i+1, price))
```

```
Predicted selling price for Client 1's home: $403,025.00
Predicted selling price for Client 2's home: $237,478.72
Predicted selling price for Client 3's home: $931,636.36
```

## Answer:

When we look at the predicted values, we see that client 3's house is costing \$931,636.36 and which seems to be justified considering the number of rooms (8) and the neighborhood poverty levels, which are much higher than the other 2 client's home.

Client 1's home is second costly and this can be inferred from the fact that a house with 5 rooms will cost less than house with 9 rooms and more than one with 4 rooms. the difference between client 1's house and client 2's house can be considered a bit absurd because 1 room cannot make a difference of  $\sim 200,000$ . There should be some other factor like location (which needs to be considered for a good model) or many others to properly justify the predictions.

Client 2's home is least costly with only 4 rooms and is justified because of high poverty percentage in neighborhood and also the no of rooms.

## Sensitivity

An optimal model is not necessarily a robust model. Sometimes, a model is either too complex or too simple to sufficiently generalize to new data. Sometimes, a model could use a learning algorithm that is not appropriate for the structure of the data given. Other times, the data itself could be too noisy or contain too few samples to allow a model to adequately capture the target variable — i.e., the model is underfitted.

**Run the code cell below to run the `fit_model` function ten times with different training and testing sets to see how the prediction for a specific client changes with respect to the data it's trained on.**

```
In [16]: vs.PredictTrials(features, prices, fit_model, client_data)
```

```
Trial 1: $391,183.33  
Trial 2: $424,935.00  
Trial 3: $415,800.00  
Trial 4: $420,622.22  
Trial 5: $418,377.27  
Trial 6: $411,931.58  
Trial 7: $399,663.16  
Trial 8: $407,232.00  
Trial 9: $351,577.61  
Trial 10: $413,700.00
```

```
Range in prices: $73,357.39
```

## Question 11 - Applicability

- In a few sentences, discuss whether the constructed model should or should not be used in a real-world setting.

**Hint:** Take a look at the range in prices as calculated in the code snippet above. Some questions to answering:

- How relevant today is data that was collected from 1978? How important is inflation?
- Are the features present in the data sufficient to describe a home? Do you think factors like quality of appliances in the home, square feet of the plot area, presence of pool or not etc should factor in?
- Is the model robust enough to make consistent predictions?
- Would data collected in an urban city like Boston be applicable in a rural city?
- Is it fair to judge the price of an individual home based on the characteristics of the entire neighborhood?

### Answer:

The constructed model cannot be used in a real-world setting because of the reasons listed below:

The data was collected back in 1978, inspite of making changes to the data for inflation, we might have not matched the real scenario. The features present in the data were relevant back in the times because of the lack of technology or facilities, like STRATIO is something that is not relevant in modern times since both the school and parents are willing to provide facilities to compensate for poor education in terms of transportation considering the abundance of facilities in today's world.

Yes some more features like quality of appliances, square feet of the plot area and other facilities can help make the data and model more relevant.

Yes the model is robust enough to make consistent predictions, considering the range of outputs but not for accurate predictions.

The data of urban area cannot be used to make a generalised model for both rural and urban areas, since the needs for rural area would be different than that of urban area and accordingly our data attributes should change.

**Note:** Once you have completed all of the code implementations and successfully answered each question above, you may finalize your work by exporting the iPython Notebook as an HTML document. You can do this by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.